

CS411

Database Systems

05: Relational Schema Design

Why Do We Learn This?

Motivation

- We have designed ER diagram, and translated it into a relational db schema $R = \text{set of } R_1, R_2, \dots$
- Now what?
- We can do the following
 - specify all relevant constraints over R
 - implement R in SQL
 - start using it, making sure the constraints always remain valid
- However, R may not be well-designed, thus causing us a lot of problems

Q: Is this a good design?

Persons with several phones:

Address	SSN	Phone Number
10 Green	123-321-99	(201) 555-1234
10 Green	123-321-99	(206) 572-4312
431 Purple	909-438-44	(908) 464-0028
431 Purple	909-438-44	(212) 555-4000

Potential Problems

- Redundancy
- Update anomalies
- Deletion anomalies

Better Designs Exist

Break the relation into two:

SSN	Address
123-321-99	10 Green
909-438-44	431 Purple

SSN	Phone Number
123-321-99	(201) 555-1234
123-321-99	(206) 572-4312
909-438-44	(908) 464-0028
909-438-44	(212) 555-4000

How do We Obtain a Good Design?

- Start with the original db schema R
- Transform it until we get a good design R^*
- Some desirable properties for R^*
 - must preserve the information of R
 - must have minimal amount of redundancy
 - must be “dependency preserving” (we’ll come to this later)
 - (must also give good query performance)

How do We Obtain a Good Design?

Normal Forms

- DB gurus have developed many “normal forms”
- Most important ones
 - Boyce-Codd, 3rd, and 4th normal forms
- If R^* is in one of these forms, then R^* is guaranteed to achieve certain good properties
 - e.g., if R^* is in Boyce-Codd NF, it is guaranteed to not have certain types of redundancy
- DB gurus have also developed algorithms to transform R into R^* that is in some of these normal forms

Normal Forms (cont.)

- DB gurus have also discussed trade-offs among normal forms
- Thus, all we have to do is
 - learn these forms
 - transform R into R^* in one of these forms
 - carefully evaluate the trade-offs
- To understand these normal forms we'll need to understand certain types of constraints
 - functional dependencies and keys

Functional Dependencies and Keys

Functional Dependencies

- A form of constraint (hence, part of the schema)
- Finding them is part of the database design
- Used heavily in schema refinement

Definition:

If two tuples agree on the attributes

$A_1, A_2 \dots A_n$

then they must also agree on the attributes

$B_1, B_2 \dots B_m$

Formally: $A_1, A_2 \dots A_n \longrightarrow B_1, B_2 \dots B_m$

Examples

EmpID	Name	Phone	Position
E0045	Smith	1234	Clerk
E1847	John	9876	Salesrep
E1111	Smith	9876	Salesrep
E9999	Mary	1234	Lawyer

- EmpID \longrightarrow Name, Phone, Position
- Position \longrightarrow Phone
- but Phone $\not\longrightarrow$ Position

More examples

Product: name, manufacturer \rightarrow price
Person: ssn \rightarrow name, age
Company: name \rightarrow stock price, president

Q: From this, can you conclude phone \rightarrow SSN?

SSN	Phone Number
123-321-99	(201) 555-1234
123-321-99	(206) 572-4312
909-438-44	(908) 464-0028
909-438-44	(212) 555-4000

Keys are a type of FD

- After defining FDs, we can now define keys
- Key of a relation R is a set of attributes that
 - functionally determines all attributes of R
 - none of its subsets determines all attributes of R
- Superkey
 - a set of attributes that contains a key

Rules about Functional Dependencies

The Splitting/Combining Rule

- $A_1A_2\dots A_n \Rightarrow B_1B_2\dots B_m$
- Equivalent to:
 - $A_1A_2\dots A_n \Rightarrow B_1;$
 - $A_1A_2\dots A_n \Rightarrow B_2;$
 - \dots
 - $A_1A_2\dots A_n \Rightarrow B_m$
- Can replace one for the other.

Trivial Functional Dependencies

- $A_1A_2\dots A_n \twoheadrightarrow A_1$
- In general,
 $A_1A_2\dots A_n \twoheadrightarrow B_1B_2\dots B_m$
if $\{B_1B_2\dots B_m\} \subseteq \{A_1A_2\dots A_n\}$
- Always holds.

Closure of a Set of Attributes

Given a set of attributes $\{A_1, \dots, A_n\}$ and a set of dependencies S .

Problem: find all attributes B such that:

any relation which satisfies S also satisfies:

$$A_1, \dots, A_n \rightarrow B$$

The **closure** of $\{A_1, \dots, A_n\}$, denoted $\{A_1, \dots, A_n\}^+$, is the set of all such attributes B

We will discuss the motivations for attribute closures soon

Algorithm to Compute Closure

Split the FDs in S so that every FD has a single attribute on the right.

Start with $X = \{A_1 A_2 \dots A_n\}$.

Repeat until X doesn't change do:

If $(B_1 B_2 \dots B_m \rightarrow C)$ is in S , such that B_1, B_2, \dots, B_m are in X and C is not in X :
add C to X .

X is now the correct value of $\{A_1 A_2 \dots A_n\}^+$

Example

A B \longrightarrow C
A D \longrightarrow E
 B \longrightarrow D
A F \longrightarrow B

Closure of {A,B}: $X = \{A, B, C, D, E\}$

Closure of {A, F}: $X = \{A, F, B, D, C, E\}$

Is this algorithm correct?

- Yes. See Text (Section 3.2.5) for proof.

Usage for Attribute Closure

- Test if X is a superkey
 - compute X^+ , and check if X^+ contains all attrs of R
- Check if $X \rightarrow Y$ holds
 - by checking if Y is contained in X^+

Reviewing Quiz 1

- Show that each of the following are not valid rules about FD's, by giving example relations that satisfy the given FDs (following the "If"), but not the FD that allegedly follows (after the "then").
 - (1) If $A \twoheadrightarrow B$ then $B \twoheadrightarrow A$
 - (2) If $AB \twoheadrightarrow C$ and $A \twoheadrightarrow C$ then $B \twoheadrightarrow C$.
- (1) $A = \text{SSN}, B = \text{Name}$
- (2) $A = \text{SSN}, B = \text{Phone}, C = \text{Name}$

Closure of a set of FDs

- Given a relation schema R & a set S of FDs
 - is the FD f logically implied by S ?
- Example
 - $R = \{A, B, C, G, H, I\}$
 - $S = A \twoheadrightarrow B, A \twoheadrightarrow C, CG \twoheadrightarrow H, CG \twoheadrightarrow I, B \twoheadrightarrow H$
 - would $A \twoheadrightarrow H$ be logically implied?
 - yes (you can prove this, using the definition of FD)
- Closure of S : $S^+ =$ all FDs logically implied by S
- How to compute S^+ ?

Computing S^+

- To check if $A \rightarrow B$ is true, we can compute A^+
- To compute all $A \rightarrow B$ implied by S , i.e., to compute the closure of S , we can use a particular algorithm.
- This algorithm depends on the so-called “Armstrong’s axioms”.

Armstrong's Axioms

- Reflexivity rule
 - $A_1A_2\dots A_n \rightarrow$ a subset of $A_1A_2\dots A_n$
- Augmentation rule
 - $A_1A_2\dots A_n \rightarrow B_1B_2\dots B_m$, then
 $A_1A_2\dots A_n C_1C_2\dots C_k \rightarrow B_1B_2\dots B_m C_1C_2\dots C_k$
- Transitivity rule
 - $A_1A_2\dots A_n \rightarrow B_1B_2\dots B_m$ and
 $B_1B_2\dots B_m \rightarrow C_1C_2\dots C_k$, then
 $A_1A_2\dots A_n \rightarrow C_1C_2\dots C_k$

Inferring S^+ using Armstrong's Axioms

- $S^+ = S$
- Loop
 - For each f in S^+ , apply reflexivity and augmentn rules
 - add the new FDs to S^+
 - For each pair of FDs in S^+ , apply the transitivity rule
 - add the new FD to S^+
- Until S^+ does not change any further

Too many rules? Not really.

- The “combining rule” can be derived from Armstrong’s axioms
 - $X \twoheadrightarrow Y$ and $X \twoheadrightarrow Z$, then $X \twoheadrightarrow YZ$
 - (X, Y, Z are sets of attributes)
- The “splitting rule” can also be derived from Armstrong’s axioms
 - $X \twoheadrightarrow YZ$, then $X \twoheadrightarrow Y$ and $X \twoheadrightarrow Z$

More on FDs

- Given a set of FDs: S .
- Let's say the set S' is *equivalent* to S , in the sense that S' can be inferred from S and vice versa.
- Any such S' is said to be a *basis* for S .
- “Minimal basis”: see Section 3.2.7.

But we were talking about schema design

- Armed with the concepts of “functional dependency” and “superkey” (and tools to reason about them), we will now define what a “good schema” is.

Desirable Properties of Schema Refinement

- 1) minimize redundancy
- 2) avoid info loss
- 3) preserve dependency
- 4) ensure good query performance

Normal Forms

First Normal Form = all attributes are atomic

Second Normal Form (2NF) = old and obsolete

Boyce Codd Normal Form (BCNF) 

Third Normal Form (3NF)

Fourth Normal Form (4NF)

Others...

Boyce-Codd Normal Form

A simple condition for removing anomalies from relations:

A relation R is in BCNF if and only if:

Whenever there is a nontrivial FD $A_1A_2\dots A_n \rightarrow B$
for R , it is the case that $A_1A_2\dots A_n$
is a super-key for R .

In English (though a bit vague):

Whenever a set of attributes of R is determining another attribute,
it should determine *all* attributes of R .

Example

Name	SSN	Phone Number
Fred	123-321-99	(201) 555-1234
Fred	123-321-99	(206) 572-4312
Joe	909-438-44	(908) 464-0028
Joe	909-438-44	(212) 555-4000

What are the dependencies?

SSN → Name

Is the left side a superkey?

Is it in BCNF?

Decompose it into BCNF

SSN	Name
123-321-99	Fred
909-438-44	Joe

SSN → Name

SSN	Phone Number
123-321-99	(201) 555-1234
123-321-99	(206) 572-4312
909-438-44	(908) 464-0028
909-438-44	(212) 555-4000


BCNF Decomposition

Find a dependency that violates the BCNF condition:

$$A_1, A_2 \dots A_n \longrightarrow B_1, B_2 \dots B_m$$

Heuristic : choose $B_1, B_2 \dots B_m$ “as large as possible”

Decompose:



Continue until
there are no
BCNF violations
left.

Example Decomposition

Person:

Name	SSN	Age	EyeColor	PhoneNumber

Functional dependencies:

$SSN \rightarrow Name, Age, Eye\ Color$

BCNF: Person1(SSN, Name, Age, EyeColor),
Person2(SSN, PhoneNumber)

What if we also had an attribute Draft-worthy, and the FD:

$Age \rightarrow Draft\text{-}worthy$

Example

- Person (Name, SSN, Age, EyeColor, Phone, Draftworthy)
- FD: SSN \rightarrow Name, Age, EyeColor
- FD: Age \rightarrow Draftworthy

Example

- Movie (title, yr, length, genre, studioName, starName)
- (Title, year, starName) is a key
- FD: Title, year \rightarrow length, genre, studioName

Example

- Movie (title, yr, studioName, President, PresAddr)
- FD: Title, yr \rightarrow studioName
- FD: studioName \rightarrow President
- FD: President \rightarrow PresAddr

Two-attribute relations

- Let A and B be the only two attributes of R
- Claim: R is in BCNF. (See Example 3.17.)
- If $A \rightarrow B$ is true, $B \rightarrow A$ is not:
- If $B \rightarrow A$ is true, $A \rightarrow B$ is not:
- If $A \rightarrow B$ is true, $B \rightarrow A$ is true:

BCNF Decomposition: The Algorithm

- Input: relation R, set S of FDs over R
 - 1) Check if R is in BCNF, if not:
 - a) pick a violation FD f: $A \rightarrow B$
 - b) compute A^+
 - c) create $R1 = A^+$, $R2 = A$ union $(R - A^+)$
 - d) compute all FDs over $R1$, using R and S. Repeat similarly for $R2$. (See Algorithm 3.12)
 - e) Repeat Step 1 for $R1$ and $R2$
 - 4) Stop when all relations are BCNF, or are two-attributes

Q: Is BCNF Decomposition unique?

Q: Does BCNF Decompos. always exist?

Properties of BCNF

- BCNF removes certain types of redundancy
 - those caused by adding many-many or one-many relationships
- For examples of redundancy that it cannot remove, see "multi-valued redundancy"
 - But this is not in curriculum

Properties of BCNF

- BCNF Decomposition avoids information loss
 - You can construct the original relation instance from the decomposed relations' instances.

Desirable Properties of Schema Refinement (again)

- 1) minimize redundancy
- 2) avoid info loss
- 3) preserve dependency
- 4) ensure good query performance


An easy decomposition?

- We saw that two-attribute relations are in BCNF.
- Why don't we break any $R(A,B,C,D,E)$ into $R_1(A,B)$; $R_2(B,C)$; $R_3(C,D)$; $R_4(D,E)$? Why bother with finding BCNF violations etc.?

Example of the “easy decomposition”

- $R = (A,B,C)$; decomposed into $R1(A,B)$; $R2(B,C)$

A	B	C
1	2	3
4	5	6



Example of the “easy decomposition”

- $R = (A,B,C)$; decomposed into $R1(A,B)$; $R2(B,C)$

A	B	C
1	2	3
4	5	6

Join

A	B	C
1	2	3
4	2	5
1	2	5
4	2	3

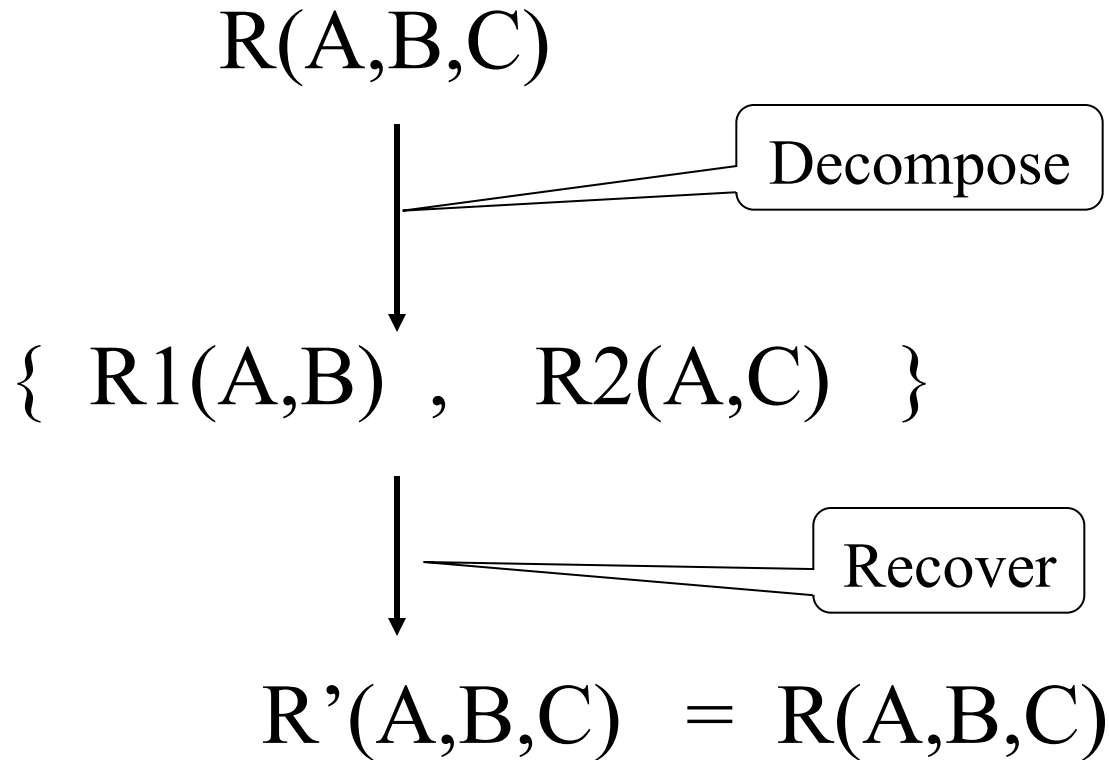
A	B
1	2
4	2

B	C
2	3
2	5

- We get back some “bogus tuples” !

Lossless Decompositions

A decomposition is *lossless* if we can recover:



R' is in general larger than R . Must ensure $R' = R$

BCNF Decomposition is Lossless

$R(A, B, C), \quad A \rightarrow C$

BCNF: $R_1(A, B), \quad R_2(A, C)$

Some tuple (a, b, c) in R
decomposes into (a, b) in R_1
and (a, c) in R_2

Recover tuples in R : $(a, b, c),$

BCNF Decomposition is Lossless

$R(A, B, C), \quad A \rightarrow C$

BCNF: $R_1(A, B), \quad R_2(A, C)$

Some tuple (a, b, c) in R	(a, b', c') also in R
decomposes into (a, b) in R_1	(a, b') also in R_1
and (a, c) in R_2	(a, c') also in R_2

Recover tuples in R : $(a, b, c),$

BCNF Decomposition is Lossless

$R(A, B, C), \quad A \rightarrow C$

BCNF: $R_1(A, B), \quad R_2(A, C)$

Some tuple (a, b, c) in R	(a, b', c') also in R
decomposes into (a, b) in R_1	(a, b') also in R_1
and (a, c) in R_2	(a, c') also in R_2

Recover tuples in R : $(a, b, c), \quad (a, b, c'), (a, b', c), (a, b', c')$ also in R

Is any of these a “bogus tuple” (not present in R)?

BCNF Decomposition is Lossless

$R(A, B, C), \quad A \rightarrow C$

BCNF: $R_1(A, B), \quad R_2(A, C)$

Some tuple (a, b, c) in R	(a, b', c') also in R
decomposes into (a, b) in R_1	(a, b') also in R_1
and (a, c) in R_2	(a, c') also in R_2

Recover tuples in R : $(a, b, c), \quad (a, b, c'), (a, b', c), (a, b', c')$ also in R

Is any of these a “bogus tuple” (not present in R)?

No! Also see text 3.4.1 for proof. (But skip 3.4.2 – 3.4.3.)

Desirable Properties of Schema Refinement (again)

- 1) **minimize redundancy**
 - 2) **avoid info loss**
 - 3) **preserve dependency**
 - 4) **ensure good query performance**
- 
- BCNF

However,

- BCNF is not always dependency preserving
- In fact, some times we cannot find a BCNF decomposition that is dependency preserving
- Can handle this situation using 3NF
- But what is “dependency preserving”?

Normal Forms


First Normal Form = all attributes are atomic

Second Normal Form (2NF) = old and obsolete

Boyce Codd Normal Form (BCNF)

Third Normal Form (3NF)

Fourth Normal Form (4NF)



Others...

3NF: A Problem with BCNF

Phone	Address	Name

FD's: $\text{Phone} \rightarrow \text{Address}$; $\text{Address, Name} \rightarrow \text{Phone}$

So, there is a BCNF violation ($\text{Phone} \twoheadrightarrow \text{Address}$), and we decompose.

Phone	Address

$\text{Phone} \rightarrow \text{Address}$

Phone	Name

No FDs

So where's the problem?

Phone	Address	Phone	Name
1234	10 Downing	1234	John
5678	10 Downing	5678	John

No problem so far. All *local* FD's are satisfied.

Let's put all the data into a single table:

Phone	Address	Name
1234	10 Downing	John
5678	10 Downing	John

Violates the dependency: Address, Name → Phone

Preserving FDs

- What if, when a relation is decomposed, the X of an $X \rightarrow Y$ ends up only in one of the new relations and the Y ends up only in another?
- Such a decomposition is not “dependency-preserving.”
- Sometimes it is not possible to decompose a relation into BCNF relations that have both lossless-join and dependency preservation properties. May need to make a tradeoff.

An alternative: 3rd Normal Form (3NF)

A simple condition for removing anomalies from relations:

A relation R is in 3rd normal form if :

Whenever there is a nontrivial dependency $A_1, A_2, \dots, A_n \rightarrow B$ for R , then $\{A_1, A_2, \dots, A_n\}$ is a super-key for R,
or B is part of a key.

“Excuse” the “Phone \rightarrow Address” FD from causing a decomposition


3NF vs. BCNF

- R is in **BCNF** if whenever $X \rightarrow A$ holds, then X is a superkey.
- Slightly stronger than 3NF.
- Example: R(A,B,C) with $\{A,B\} \rightarrow C$, $C \rightarrow A$
 - 3NF but not BCNF

Decomposing R into 3NF

- The algorithm is complicated
- 1. Get a “minimal basis” G of given FDs (Section 3.2.7)
- 2. For each FD $X \rightarrow A$ in the minimal basis G , use XA as the schema of a new relation.
- 3. If none of the schemas from Step 2 is a superkey, add another relation whose schema is a key for the original relation.
- Result will be lossless, will be dependency-preserving, 3NF; might not be BCNF
- *Example 3.27 in textbook.*
- But skip Section 3.5.3.

Desirable Properties of Schema Refinement (again)

- 1) minimize redundancy
 - 2) **avoid info loss**
 - 3) **preserve dependency**
 - 4) ensure good query performance
- 

Fact of life...

Finding a decomposition which is both lossless and dependency-preserving is not always possible.

Guideline: Aim for BCNF and settle for 3NF

Multi-valued Dependencies and 4NF

we will not cover this.

Caveat

- Normalization is not the be-all and end-all of DB design
- Example: suppose attributes A and B are always used together, but normalization theory says they should be in different tables.
 - decomposition might produce unacceptable performance loss (extra disk reads)