

# *Introduction to Database Programming*

Presented by: Yunliang Jiang  
2010.10.14

# DB Programming

- SQL is a very high-level language.
- **Not** intended for general-purpose computations.
- Solutions:
  - Outside DBMS: use SQL together with general-purpose programming languages.
  - Inside DBMS: augment SQL with constructs from general-purpose programming languages.

# All these methods follow the same basic paradigm

1. Connect to a DB server.
2. Say what database you want to use.
3. Assemble a string containing an SQL statement.
4. Get the DBMS to prepare a plan for executing the statement.
5. Execute the statement.
6. Extract the results into variables in the local programming language.

# Overview

- Outside

- API Approach:

- Vendor specific libraries [80's- ]

- MySQL API for PHP

- Open interfaces [90's - ]

- JDBC, ODBC

- Embedded SQL [70's- ]

- Embedded SQL for C/C++.

- Not widely used.

- Inside

- Stored procedures/functions: [80's- ]

# API Approach

- Programmer uses a library of functions or classes and call them as part of an ordinary C or Java program.
- SQL commands are sent to the DBMS at runtime.
- These API's are based on an standard called SQL/CLI = “Call-Level Interface.”

# Database specific API

- Designed and implemented by the DBMS vendor for a specific programming language.
- We go over MySQL API for PHP.

# Connection

- `mysql_connect` opens a connection to the DBMS.
- It gets the DBMS and login information and returns a *connection* resource. The *connection* resource is used in future calls.
- `mysql_select_db` selects the desired database
- `mysql_close` closes the connection at the end. It is called automatically at the end of the script.

# Executing Statements

- `mysql_query(S, C)` causes the SQL statement represented by `S` to be executed on connection `C`.
- We can detect DBMS initiated errors using `mysql_error()`.
- `mysql_query` returns a handle to the query result set.
  - It returns `TRUE/FALSE` for `DELETE`, `UPDATE`, and `INSERT` statements.



# Fetching Tuples

- When the SQL statement executed is a query, we need to fetch the tuples of the result.
- `mysql_fetch_array(H)` gets the tuples from the result set *H*.
- `mysql_free_result(H)` frees the result set. It is called automatically at the end of the script

# Fetching Tuples

- Random access fetching by `mysql_data_seek`.
- More functions at:
  - <http://us.php.net/mysql>

# Open interface: JDBC

- Database specific API makes the program dependent to one DBMS.
- Open interfaces solve this problem:
  - Designed by a third party like the creator of the language.
  - Implemented by DBMS vendors.
  - Used by DB programmers.
- Programmer do not have to change their code to work with another DBMS.
- <http://www.jdbc-tutorial.com/>

# Connections and Statements

- The same progression from connections to statements that we saw in CLI appears in JDBC.
- A *connection object* is obtained from the environment in a somewhat implementation-dependent way.
- We'll start by assuming we have myCon, a connection object.

# Statements

- JDBC provides two classes:
  1. Statement = an object that can accept a string that is an SQL statement and can execute such a string.
  2. PreparedStatement = an object that has an associated SQL statement ready to execute.
- Why? Performance
  - PreparedStatement SQL command is stored in DBMS after the first call.

# Creating Statements

- The Connection class has methods to create Statements and PreparedStatement.

```
Statement stat1 = myCon.createStatement();
```

```
PreparedStatement stat2 =
```

```
myCon.prepareStatement(
```

```
    "SELECT Address FROM Supplier" +
```

```
    "WHERE Supplier_Name = 'John Smith'"
```

```
);
```

Java trick: +  
concatenates  
strings.



# Executing SQL Statements

- Programmer handles errors using `SQLException` class.
- JDBC distinguishes queries from modifications, which it calls “updates.”
- `Statement` and `PreparedStatement` each have methods `executeQuery` and `executeUpdate`.
  - For `Statements`, these methods have one argument: the query or modification to be executed.
  - For `PreparedStatement`s: no argument.

# Example: Update

- stat1 is a Statement.
- We can use it to insert a tuple as:

```
stat1.executeUpdate(  
    “INSERT INTO Supplier” +  
    “VALUES(‘S4’, ‘Mary’, ‘12 Goodwin St.’)”  
);
```



## Example: Query

- stat2 is a PreparedStatement holding the query “SELECT Address FROM Supplier WHERE Supplier\_Name = ‘John Smith’”.
- executeQuery returns an object of class ResultSet --- we’ll examine it later.
- The query:

```
ResultSet Menu = stat2.executeQuery();
```

# Accessing the ResultSet

- An object of type ResultSet is something like a cursor.
- Method Next() advances the “cursor” to the next tuple.
  - The first time Next() is applied, it gets the first tuple.
  - If there are no more tuples, Next() returns the value FALSE.

# Accessing Components of Tuples

- When a ResultSet is referring to a tuple, we can get the components of that tuple by applying certain methods to the ResultSet.
- Method `getX (i )`, where  $X$  is some type, and  $i$  is the component number, returns the value of that component.
  - The value must have type  $X$ .

# Example: Accessing Components

- List is the ResultSet for the query “SELECT Address FROM Supplier WHERE Supplier\_Name = ‘John Smith’”.

- Access the address from each tuple by:

```
while ( List.Next() ) {  
    theAddress = List.getString(1);  
    /* do something */  
}
```

- More: <http://java.sun.com/products/jdbc/overview.html>

# Programming Inside the DBMS

# Stored Procedures

- An extension to SQL, called SQL/PSM, or “persistent, stored modules,” allows us to store procedures as database schema elements.
- The programming style is a mixture of conventional statements (if, while, etc.) and SQL.
- Lets us do things we cannot do in SQL alone.
- Why? Performance.
- They are harder to develop and maintain.

# Basic PSM Format

```
CREATE PROCEDURE <name> (  
    <parameter list> )  
    <optional local declarations>  
    <body>;
```

- Function alternative:

```
CREATE FUNCTION <name> (  
    <parameter list> ) RETURNS <type>
```

Example: compute square footage of house lot; then you can query on it

# Parameters in PSM

- Unlike the usual name-type pairs in languages like C, PSM uses mode-name-type triples, where the *mode* can be:
  - IN = procedure uses value, does not change value.
  - OUT = procedure changes, does not use.
  - INOUT = both.



# Example: Stored Procedure

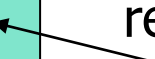
- Let's write a procedure that takes three arguments  $i$ ,  $n$ , and  $a$ , and adds a tuple to Supplier that has Supplier\_ID =  $i$ , Supplier\_Name =  $n$ , and Address =  $a$ .
  - Used to add new Supplier more easily.

# The Procedure

```
CREATE PROCEDURE newSup (
```

```
IN i    VARCHAR(20),  
IN n    VARCHAR(50),  
IN a    VARCHAR(100)
```

Parameters are all  
read-only, not changed



```
)
```

```
INSERT INTO Supplier  
VALUES(i, n, a);
```

The body ---  
a single insertion



# Invoking Procedures

- Use SQL/PSM statement CALL, with the name of the desired procedure and arguments.

- Example:

CALL newSup('s5', 'Bob', '1 Main St.);

- Functions used in SQL expressions where a value of their return type is appropriate.

# Types of PSM statements -- 1

- RETURN <expression> sets the return value of a function.
  - Unlike C, etc., RETURN *does not* terminate function execution.
- DECLARE <name> <type> used to declare local variables.
- BEGIN . . . END for groups of statements.
  - Separate by semicolons.

# Types of PSM Statements -- 2

- Assignment statements:

SET <variable> = <expression>;

– Example: SET b = 'Bud';

- Statement labels: give a statement a label by prefixing a name and a colon.

# Example: IF

- Let's rate suppliers by how many parts they have
  - $\leq 0$  parts: 'bad'.
  - $< 2$  parts: 'average'.
  - $\geq 2$  parts: 'good'.
- Function Rate(s) rates supplier s.

# Example: IF (continued)

```
CREATE FUNCTION Rate (IN b VARCHAR(50) )
```

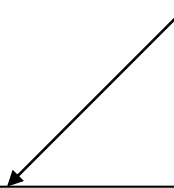
```
  RETURNS CHAR(10)
```

```
  DECLARE pCnt INTEGER;
```

```
  BEGIN
```

```
    SET pCnt = (SELECT COUNT(*) FROM Supplier,  
                Catalog WHERE Supplier.Supplier_ID =  
                Catalog.Supplier_ID and  
                Supplier_Name= b);
```

Number of  
the parts of the  
supplier b




```
    IF pCnt < 0 THEN RETURN 'bad'  
    ELSEIF pCnt < 2 THEN RETURN 'average'  
    ELSE RETURN 'good'  
    END IF;
```

Nested  
IF statement



```
  END;
```

Return occurs here, not at  
one of the RETURN statements



# Example: Exiting a Loop

loop1: LOOP

...

LEAVE loop1;      ← If this statement is executed . . .

...

END LOOP;      ← Control winds up here



# Other Loop Forms

- WHILE <condition>  
    DO <statements>  
END WHILE;
- REPEAT <statements>  
UNTIL <condition>  
END REPEAT;

# Queries

- General SELECT-FROM-WHERE queries are *not* permitted in PSM.
- There are three ways to get the effect of a query:
  1. Queries producing one value can be the expression in an assignment.
  2. Single-row SELECT . . . INTO.
  3. Cursors.

# Example: Assignment/Query

- If  $a$  is a local variable and Supplier(Supplier\_ID, Supplier\_Name, Address) the usual relation, we can get the address of 'John Smith' by:

```
SET a = (SELECT Address FROM Supplier  
        WHERE Supplier_Name = 'John  
        Smith' );
```

# SELECT . . . INTO

- An equivalent way to get the value of a query that is guaranteed to return a single tuple is by placing INTO <variable> after the SELECT clause.
- Example:

```
SELECT Address INTO a FROM  
Supplier WHERE Supplier_Name =  
'John Smith';
```

# Cursors

- A *cursor* is essentially a tuple-variable that ranges over all tuples in the result of some query.
- Declare a cursor *c* by:

DECLARE *c* CURSOR FOR <query>;

# Opening and Closing Cursors

- To use cursor  $c$ , we must issue the command:  
    OPEN  $c$ ;
  - The query of  $c$  is evaluated, and  $c$  is set to point to the first tuple of the result.
- When finished with  $c$ , issue command:  
    CLOSE  $c$ ;

# Fetching Tuples From a Cursor

- To get the next tuple from cursor  $c$ , issue command:

FETCH FROM  $c$  INTO  $x_1, x_2, \dots, x_n$  ;

- The  $x$  's are a list of variables, one for each component of the tuples referred to by  $c$ .
- $c$  is moved automatically to the next tuple.

# Breaking Cursor Loops -- 1

- The usual way to use a cursor is to create a loop with a FETCH statement, and do something with each tuple fetched.
- A tricky point is how we get out of the loop when the cursor has no more tuples to deliver.



# Breaking Cursor Loops -- 2

- Each SQL operation returns a *status*, which is a 5-digit number.
  - For example, 00000 = “Everything OK,” and 02000 = “Failed to find a tuple.”
- In PSM, we can get the value of the status in a variable called SQLSTATE.

# Breaking Cursor Loops -- 3

- We may declare a condition, which is a boolean variable that is true if and only if SQLSTATE has a particular value.
- Example: We can declare condition NotFound to represent 02000 by:

```
DECLARE NotFound CONDITION FOR  
SQLSTATE '02000';
```

# Breaking Cursor Loops -- 4

- The structure of a cursor loop is thus:

```
cursorLoop: LOOP
```

```
...
```

```
FETCH c INTO ... ;
```

```
IF NotFound THEN LEAVE cursorLoop;
```

```
END IF;
```

```
...
```

```
END LOOP;
```

## Example: Cursor

- Let's write a procedure that examines Supplier and Catalog and raises by \$1 the cost of all parts in the 'John Smith' inventory that are under \$10.
  - Yes, we could write this as an UPDATE, but the details are instructive anyway.

# The Needed Declarations

```
CREATE PROCEDURE raiseCost( )
```

```
DECLARE theID VARCHAR(50);  
DECLARE theCost REAL;
```

Used to hold  
ID-cost pairs  
when fetching  
through cursor c

```
DECLARE NotFound CONDITION FOR  
SQLSTATE '02000';
```

```
DECLARE c CURSOR FOR
```

Returns John's parts costs

```
(SELECT Supplier.Supplier_ID, Cost FROM  
Supplier, Catalog WHERE Supplier.Supplier_ID =  
Catalog.Supplier_ID and Supplier_Name = 'John  
Smith');
```

# The Procedure Body


```
BEGIN
```

```
  OPEN c;
```

```
  menuLoop: LOOP
```

```
    FETCH c INTO theID, theCost;
```

Check if the recent  
FETCH failed to  
get a tuple



```
    IF NotFound THEN LEAVE menuLoop END IF;
```

```
    IF theCost < 10.00 THEN
```

```
      UPDATE Catalog SET Cost = theCost+1.00
```

```
      WHERE Supplier_ID = theID;
```


```
    END IF;
```

```
  END LOOP;
```

```
  CLOSE c;
```

```
END;
```

If John charges less than \$10 for  
the part raise it's cost by \$1.



# Further Readings

- <http://dev.mysql.com/doc/refman/5.0/en/stored-routines.html>
- [http://www.oracle.com/technology/tech/pl\\_sql/index.html](http://www.oracle.com/technology/tech/pl_sql/index.html)

# Embedded SQL



# Embedded SQL

- A standard for combining SQL with different languages.
- Key idea: Use a preprocessor to turn SQL statements into procedure calls that fit with the host-language code surrounding.
- All embedded SQL statements begin with EXEC SQL, so the preprocessor can find them easily.

# Shared Variables

- To connect SQL and the host-language program, the two parts must share some variables.
- Declarations of shared variables are bracketed by:

Always  
needed

```
EXEC SQL BEGIN DECLARE SECTION;  
    <host-language declarations>  
EXEC SQL END DECLARE SECTION;
```

# Use of Shared Variables

- In SQL, the shared variables must be preceded by a colon.
  - They may be used as constants provided by the host-language program.
  - They may get values from SQL statements and pass those values to the host-language program.
- In the host language, shared variables behave like any other variable.

# Example: Looking Up Prices

- We'll use C with embedded SQL to sketch the important parts of a function that obtains a beer and a bar, and looks up the price of that beer at that bar.
- Assumes database has our usual Sells(bar, beer, price) relation.

# Example: C Plus SQL

```
EXEC SQL BEGIN DECLARE SECTION;
```

```
char theID[21];
```

```
float theCost;
```

Note 21-char  
arrays needed  
for 20 chars +  
endmarker

```
EXEC SQL END DECLARE SECTION;
```

```
/* obtain values for theID and theCost */
```

```
EXEC SQL SELECT Cost INTO :theCost  
FROM Catalog  
WHERE Supplier_ID = theID;
```

```
/* do something with theCost */
```

SELECT-INTO  
just like PSM

# Embedded Queries

- Embedded SQL has the same limitations as PSM regarding queries:
  - You may use SELECT-INTO for a query guaranteed to produce a single tuple.
  - Otherwise, you have to use a cursor.
    - Small syntactic differences between PSM and Embedded SQL cursors, but the key ideas are identical.

# Cursor Statements

- Declare a cursor *c* with:

EXEC SQL DECLARE *c* CURSOR FOR <query>;

- Open and close cursor *c* with:

EXEC SQL OPEN CURSOR *c*;

EXEC SQL CLOSE CURSOR *c*;

- Fetch from *c* by:

EXEC SQL FETCH *c* INTO <variable(s)>;

- Macro NOT FOUND is true if and only if the FETCH fails to find a tuple.

## Example -- 1

- Let's write C + SQL to print list of Supplier's with the name 'John Smith'.
- A cursor will visit each Supplier tuple that has Supplier\_Name = 'John Smith'.



## Example – 2 (Declarations)

```
EXEC SQL BEGIN DECLARE SECTION;
```

```
    char theID[21]; char theAddress[151];
```

```
EXEC SQL END DECLARE SECTION;
```

```
EXEC SQL DECLARE c CURSOR FOR  
    SELECT theID, theAddress FROM Supplier  
    WHERE Supplier_Name = 'John Smith';
```



The cursor declaration goes  
outside the declare-section

## Example – 3 (Executable)

```
EXEC SQL OPEN CURSOR c;
```

```
while(1) {
```

```
    EXEC SQL FETCH c
```

```
        INTO :theID, :theAddress;
```

```
    if (NOT FOUND) break;
```

```
    /* format and print theID and theAddress */
```

```
}
```

```
EXEC SQL CLOSE CURSOR c;
```

The C style  
of breaking  
loops



# Need for Dynamic SQL

- Most applications use specific queries and modification statements in their interaction with the database.
  - Thus, we can compile the EXEC SQL ... statements into specific procedure calls and produce an ordinary host-language program that uses a library.
- What if the program is something like a generic query interface, that doesn't know what it needs to do until it runs?

# Dynamic SQL

- Preparing a query:

```
EXEC SQL PREPARE <query-name>  
          FROM <text of the query>;
```

- Executing a query:

```
EXEC SQL EXECUTE <query-name>;
```

- “Prepare” = optimize query.
- Prepare once, execute many times.

# Example: A Generic Interface

```
EXEC SQL BEGIN DECLARE SECTION;
```

```
    char query[MAX_LENGTH];
```

```
EXEC SQL END DECLARE SECTION;
```

```
while(1) {
```

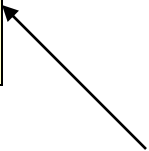
```
    /* issue SQL> prompt */
```

```
    /* read user's query into array query */
```

```
    EXEC SQL PREPARE q FROM :query;
```

```
    EXEC SQL EXECUTE q;
```

```
}
```



q is an SQL variable  
representing the optimized  
form of whatever statement  
is typed into :query

# Execute-Immediate

- If we are only going to execute the query once, we can combine the PREPARE and EXECUTE steps into one.
- Use:

**EXEC SQL EXECUTE IMMEDIATE <text>;**

# Example: Generic Interface Again

```
EXEC SQL BEGIN DECLARE SECTION;
    char query[MAX_LENGTH];
EXEC SQL END DECLARE SECTION;
while(1) {
    /* issue SQL> prompt */
    /* read user's query into array
    query */
    EXEC SQL EXECUTE IMMEDIATE :query;
}
```

# Further Readings

<http://infolab.stanford.edu/~ullman/fcdb/oracle/or-proc.html>