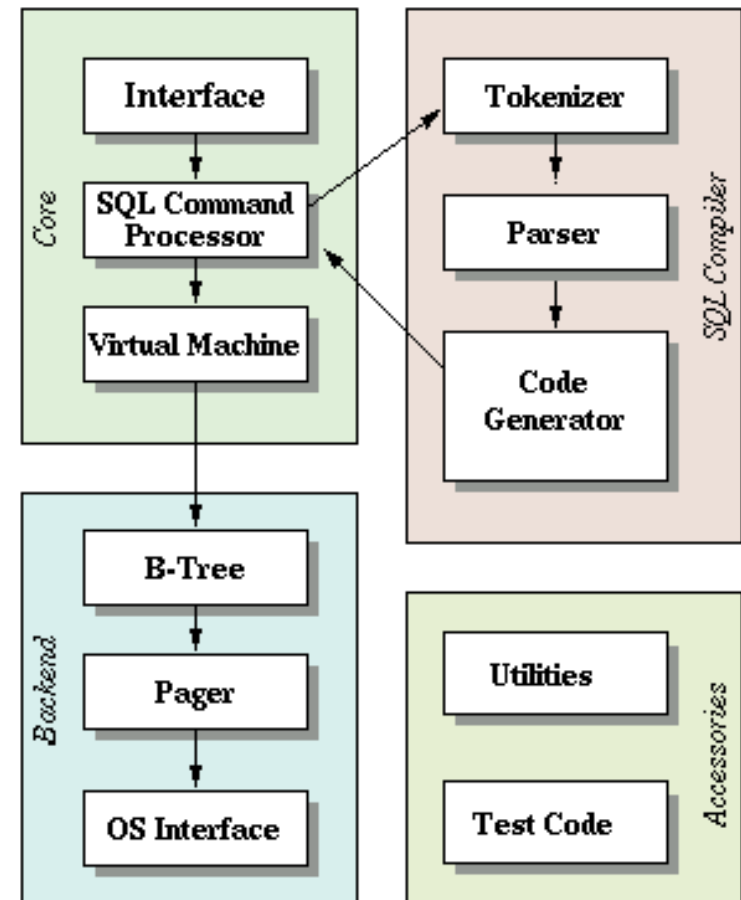


Anatomy of a Database System

Mangesh Bendre
<mangeshbendre@gmail.com>

Block Diagram of SQLite

- Execution of SQL
 1. Interface
 2. Tokenizer
 3. Parser
 4. Code Generator for VM
 5. Virtual DB Engine
 6. B-tree
 7. Pager
 8. OS interface



1. Interface

- shell.c
- Provides a command line interface to accept the SQL, execute it and display the results.
- When SQLite is used as a library, the application performs duties of an interface by calling the SQLite API
 - sqlite3_open
 - sqlite3_exec
 - sqlite3_close

1. Interface (Example)

- Execute a SQL via the API

```
char *sql = "SELECT * FROM EMPLOYEE;";
rc = sqlite3_open("my.db", &db);
rc = sqlite3_exec(db, sql, callback, 0, &zErrMsg);
sqlite3_close(db);
static int callback(void *NotUsed, int argc, char **argv, char **azColName){
    int i;
    for(i=0; i<argc; i++)
        printf("%s = %s\n", azColName[i], argv[i] ? argv[i] : "NULL");
    return 0;
}
```

2. Tokenizer

- tokenize.c
- Parse the input SQL to tokenize it.
- Tokens defined as TK_ (Tokens defined in Parser used in tokenizer)
- Tokens defined here are available in the parser.

2. Tokenizer (Example)

- The SQL is tokenized as below:

1.	TK_SELECT	SELECT
2.	TK_SPACE	
3.	TK_STAR	*
4.	TK_SPACE	
5.	TK_FROM	FROM
6.	TK_SPACE	
7.	TK_ID	EMPLOYEE
8.	TK_SEMI	;

3. Parser

- `parse.y`
- LEMON parser – LALR (Look Ahead Left to Right) Parser
- Grammar rules to parse SQL
- Tokens defined in Parser will be available in the tokenizer.

3. Parser (Example)

- The SELECT statement matches the following grammar construct:

```
cmd ::= select(X). {
```

```
    SelectDest dest = {SRT_Output, 0, 0, 0, 0};
```

```
    sqlite3Select(pParse, X, &dest);
```

```
    sqlite3SelectDelete(pParse->db, X);
```

```
}
```

```
oneselect(A) ::=
```

```
    SELECT distinct(D) selcollist(W) from(X) where_opt(Y)
```

```
    groupby_opt(P) having_opt(Q) orderby_opt(Z) limit_opt(L). {
```

```
A = sqlite3SelectNew(pParse, W, X, Y, P, Q, Z, D, L.pLimit, L.pOffset);
```

```
}
```

```
selcollist(A) ::= sclp(P) STAR.
```

```
from(A) ::= FROM seltablist(X).
```


4. Code Generator for VM

- `select.c/where.c/insert.c/expr.c`
- Handles VDBE code generation for the SQL
- Parser based on the grammar calls the respective function to generate the VDBE code.
- The VDBE code is generated based on the logic of the SQL.
- SQL Optimization is also done as part of code generation

4. Code Generator for VM (Example)

- `sqlite3Select (select.c)`
- Start Code generation with
`v = sqlite3GetVdbe (pParse) ;`
- Generate Op-codes with
`sqlite3VdbeAddOp4`

4. Code Generator for VM (Example)

```
sqlite> explain select * from employee;
```

addr	opcode	p1	p2	p3	p4	p5	comment
----	-----	----	----	----	-----	--	-----
0	Trace	0	0	0		00	
1	Goto	0	10	0		00	
2	OpenRead	0	2	0	2	00	employee
3	Rewind	0	8	0		00	
4	Column	0	0	1		00	employee.name
5	Column	0	1	2		00	employee.age
6	ResultRow	1	2	0		00	
7	Next	0	4	0		01	
8	Close	0	0	0		00	
9	Halt	0	0	0		00	
10	Transaction	0	0	0		00	
11	VerifyCookie	0	1	0		00	
12	TableLock	0	2	0	employee	00	
13	Goto	0	2	0		00	

5. Virtual DB Engine

- `vdbe.c`
- Similar to assembly programs
- Linear sequence of operations with Op-codes and corresponding Operands
- Operands
 - P1, P2 and P3 are Integers
 - P4 is a null terminated String
 - P5 is a unsigned character

5. Virtual DB Engine (cont.)

- Big SWITCH statement with different Op-codes as different case statements.
- Column P1 P2 P3 P4 P5
 - Interpret the data that cursor P1 points to as a structure built using the MakeRecord instruction.
 - Extract the P2th column from this record. If there are less than (P2+1) values in the record, extract a NULL.

Record Format: Variable Length Integer (VARINT)

- Static Huffman encoding of 64-bit two's-complement integers

Bytes	Value	Bit Pattern
1	7 bit	0xxxxxxx
2	14 bit	1xxxxxxx 0xxxxxxx
3	21 bit	1xxxxxxx 1xxxxxxx 0xxxxxxx
4	28 bit	1xxxxxxx 1xxxxxxx 1xxxxxxx 0xxxxxxx
5	35 bit	1xxxxxxx 1xxxxxxx 1xxxxxxx 1xxxxxxx 0xxxxxxx
.		
.		

Record Format: Record Structure

- The Record structure is stored for each of the record.
- The header size and the types are stored as varints.

	hdr-size		type 0		type 1		...		type N-1		data0		...		data N-1	
--	----------	--	--------	--	--------	--	-----	--	----------	--	-------	--	-----	--	----------	--

Record Format: Serial Type

- The Serial type identifies the data type. Column data type is irrelevant.

serial type	bytes of data	type
-----	-----	-----
0	0	NULL
1	1	signed integer
2	2	signed integer
3	3	signed integer
4	4	signed integer
5	6	signed integer
6	8	signed integer
7	8	IEEE float
8	0	Integer constant 0
9	0	Integer constant 1
10,11		reserved for expansion
N>=12 and even	$(N-12)/2$	BLOB
N>=13 and odd	$(N-13)/2$	text

6. B-tree

- btree.c
- A single B-Tree structure is stored using one or more database pages.
- A page contains a single B-tree node.
 - Table B-tree
 - 64-bit integer as keys
 - Index B-tree
 - Database records as keys

7. Pager

- pager.c
- Used to access a DB file.
- Caching of pages.
- It implements atomic commits/rollbacks by use of a journal file.
- Implements file locking.
 - Exclusive lock while writing.
 - Shared lock while reading.

8. OS Interface

- `os.c` – `os_unix.c/os_win.c`
- Abstraction layer to interface with the Operation System

FreeDB – A schema free DB

What is FreeDB

- Schema Free tables.
 - Values as name value pairs - Pair
(“id”, “1001”)
 - Values as Lists - List
[10, 20, 30, 40]
- By combination of the above two constructs flexible structures can be constructed.
 - [(“id”, “1001”), (“name”, “MSB”)]

Implementation of Free DB using SQLite

High level Implementation

- Tokenizer
 - To tokenize new tokens ('[' and ']')
- Parser
 - To support new SQL syntax (List & Pair) for free DB
- Code Generator
 - Generate new Opcodes to handle (List & Pair)
- VDBE
 - Execute new Opcodes

Tokenizer (tokenize.c)

- Support for List by use of [and]

```
case ']' : {
    *tokenType = TK_RB;
    return 1; }
case '[' : {
    if (comma_found) {
        *tokenType = TK_LB;
        return 1; }
    else
        *tokenType = c==']' ? TK_ID : TK_ILLEGAL;
    return i;
}
```


Parser

- Syntax for List and Pair

```
expr(A) ::= pair(X). {A = X;}
expr(A) ::= list(X). {A = X;}
/* ("price", 99.99) */
pair(A) ::= LP expr(X) COMMA expr(Y) RP.
{ spanBinaryExpr
  (&A, pParse, TK_PAIR, &X, &Y); }
/* [61820, 61821, 61822] */
list(A) ::= LB itemlist(X) RB. { A.pExpr =
sqlite3PExpr(pParse, TK_LIST, 0, 0, 0);
A.pExpr->x.pList = X; }
```

Expression Processing (expr.c) - List

- Inserts a new op-code **OP_List** and then adds the list of the expressions that form the list

```
case TK_LIST:
{
    int n;
    sqlite3VdbeAddOp2(v, OP_List, pExpr-
>x.pList->nExpr, inReg);
    for (n=0;n<pExpr->x.pList->nExpr;n++)
    {
        inReg = sqlite3ExprCodeTarget(pParse,
pExpr->x.pList->a[n].pExpr, inReg+1) ;
    }
}
```

Expression Processing (expr.c) - Pair

- Inserts a new op-code **OP_StringPair** and then adds the two expressions that form the pair.

```
case TK_PAIR: {  
    sqlite3VdbeAddOp2(v, OP_StringPair, 2, inReg);  
  
    inReg = sqlite3ExprCodeTarget(pParse, pExpr->pLeft,  
    inReg+1);  
    inReg = sqlite3ExprCodeTarget(pParse, pExpr->pRight,  
    inReg+1) ;  
}
```

VDBE (vdbec)

- Define a case statement to define a new op-code.

```
case OP_StringPair: {  
    assert( pOp->p1 == 2 );  
    pOut->u.i = pOp->p1;  
    pOut->aux_flags = MEM_Pair; break;  
}
```

Record Format Serial Type

- Two new serial types to identify the new data types

serial type	bytes of data	type
-----	-----	-----
N \geq 12 and 00	$(N-12)/4$	BLOB
N \geq 13 and 01	$(N-13)/4$	text
N \geq 12 and 10	$(N-12)/4$	Pair
N \geq 13 and 11	$(N-13)/4$	List

Thank You